

Deep Generative Models

Chapter 5: Variational Inference and VAEs

Ali Bereyhi

ali.bereyhi@utoronto.ca

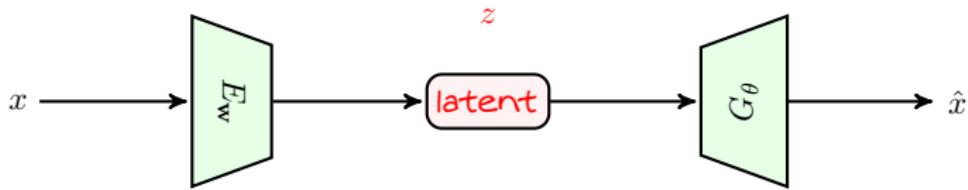
Department of Electrical and Computer Engineering
University of Toronto

Summer 2025

Early Work on VAE

For some time: people considered using *decoder* of AEs to generate data

- ① train an encoder-decoder architecture
- ② sample *latent* at random and pass it through decoder



The idea however does not work in this basic form, because

prior distribution of the *latent* is unknown

Early Work: Kingma and Welling

Kingma and Welling used variational inference to *control the latent prior*
this somehow gave birth to VAEs

In their basic experiments, they considered MLP *encoder and decoder*

- *decoder computes feature $h = \tanh(\text{Linear}(z))$ and then*

$$\mu = \text{Linear}(h)$$

$$\log \sigma = \text{Linear}(h)$$

- *encoder works exactly the same with x and z swapped*

Early Work: Kingma and Welling

The results on *Frey Face dataset* is acceptable



and the impact of *latent size* can be easily seen via *MNIST*



(a) 2-D latent space

(b) 5-D latent space

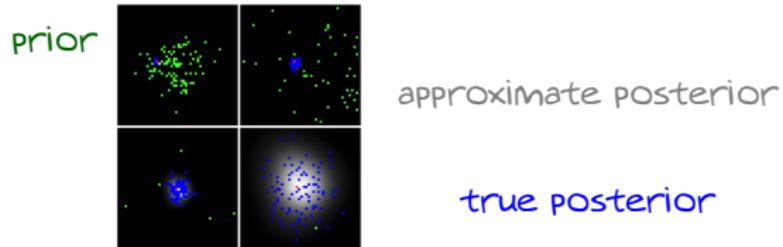
(c) 10-D latent space

(d) 20-D latent space

Early Work: Rezende, Mohamed, and Wierstra

Kingma and Welling motivated the others to dive deeper

- Rezende et al. provided solid base for VAE
- They looked into the stochastic computation graph of VAEs
- They visualized the **true** and approximate posterior

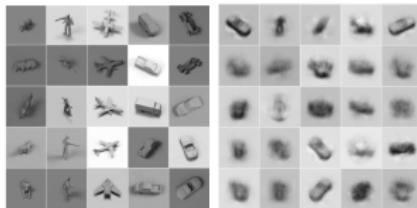


In their experiment, they considered simple MLP encoder and decoder

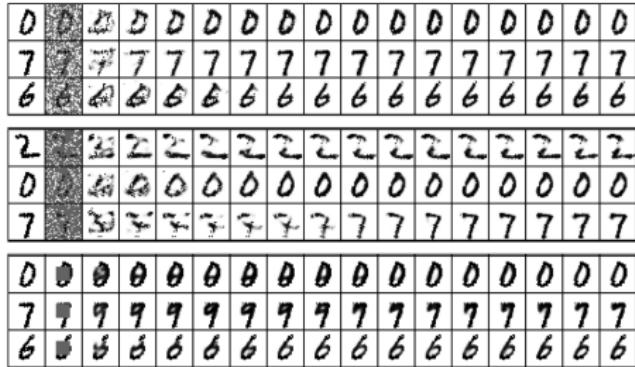
- both modules were **three-layer** MLPs
- they also investigated other applications, e.g., image inpainting

Early Work: Rezende, Mohamed, and Wierstra

Since model is simple, we get lower quality on more **complex** datasets like NORB



But on **simple MNIST**, we can do very well



Deep Convolutional VAE

People started to use CNNs to better encode and decode

4	0	1	2	3	4	5	6	7	8	9
9	0	1	2	3	4	5	6	7	8	9
5	0	1	2	3	4	5	6	7	8	9
4	0	1	2	3	4	5	6	7	8	9
2	0	1	2	3	4	5	6	7	8	9
7	0	1	2	3	4	5	6	7	8	9
5	0	1	2	3	4	5	6	7	8	9
1	0	1	2	3	4	5	6	7	8	9
7	0	1	2	3	4	5	6	7	8	9
1	0	1	2	3	4	5	6	7	8	9

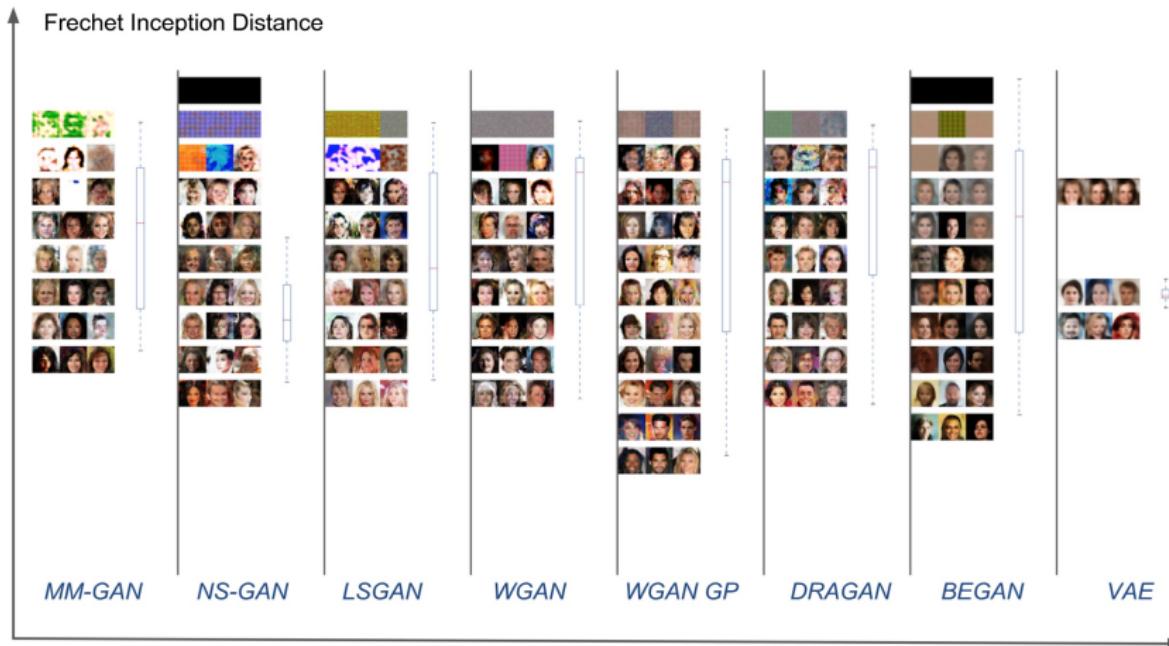
(b) MNIST analogies

10	11	12	13	14	15	16	17	18	19	10
15	11	12	13	14	15	16	17	18	19	10
36	10	20	30	40	50	60	70	80	90	00
7	1	2	3	4	5	6	7	8	9	10
13	11	12	13	14	15	16	17	18	19	10
30	11	12	13	14	15	16	17	18	19	10
61	11	21	31	41	51	61	71	81	91	01
20	10	20	30	40	50	60	70	80	90	00
28	11	22	33	44	55	26	27	28	29	20
22	11	22	33	44	55	26	27	28	29	20

(c) SVHN analogies

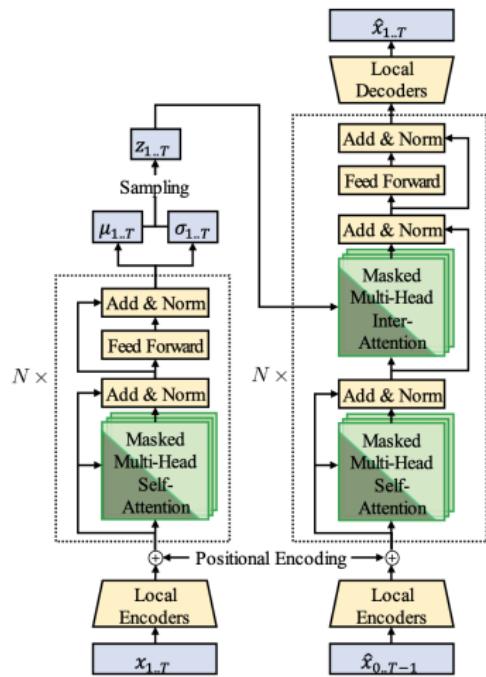
ResNet-VAEs

The results motivated people to further improve encoder and decoder



Transformer-VAEs

The results motivated people to further improve encoder and decoder



VAE Limits: Sharpness

Despite computational advances: *there are a few key restrictions to VAEs*

- ① VAE visual output is quite *blurry* ↵ for a generic data: *not too sharp*
 - ↳ This is observed simply by comparing VAE and GAN outputs



- ↳ The reason is *intuitively understandable*
 - The output is sampled conditionally from a *Gaussian distribution*
 - This means that after generating $\mu = G_\theta(z)$ we compute

$$x = \mu + \varepsilon \sim \mathcal{N}(0, \sigma^2)$$

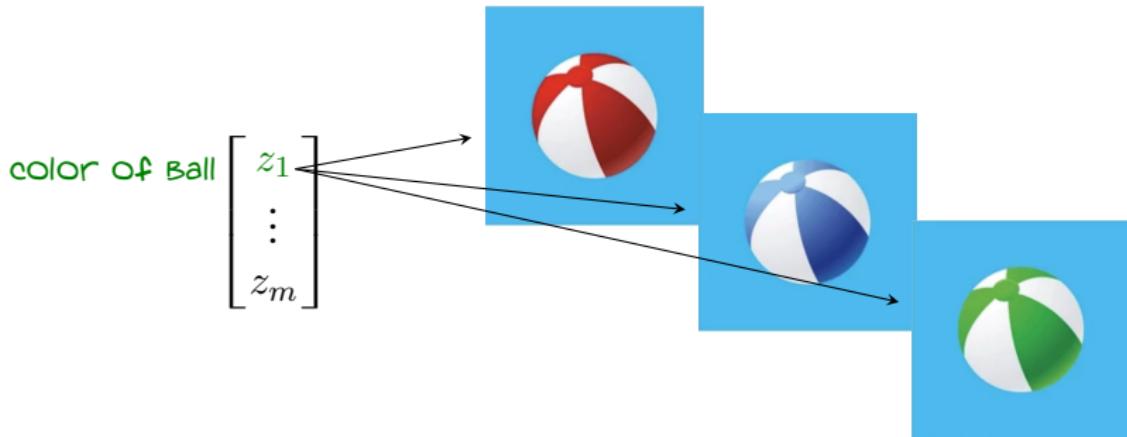
which is always a *noisy image!*

Disentanglement

Disentanglement in Latent Representation

A **latent** representation $z \in \mathbb{R}^m$ is **well-disentangled** if

- each z_i corresponds to a **single** generative factor, e.g., position or color
- changing z_i **only** affects the **corresponding factor** leaving others invariant



Disentanglement via β -VAE

In VAEs: we compute a **latent** representation

- If latent is **well-disentangled** we can **better modify generation**
- This is an **ideal property** for the decoder

② VAEs have **limited** control over latent **disentanglement**

- ↳ There is **no** parameter in encoder or decoder to impose such property
- ↳ Intuitively, we can improve **disentanglement** by **increasing latent dimension**
- ↳ Larger latent size however **introduces redundant features** as well

A better approach is the so-called β -VAE modification

$$\beta - \text{ELBO}(\mathbf{w}, \theta | \mathbf{x}) = \hat{\mathbb{E}}_{z \sim Q_{\mathbf{w}}} \{ \log P_{\theta}(\mathbf{x}|z) \} - \beta D_{\text{KL}}(Q_{\mathbf{w}} \| \mathcal{N}^0)$$

- ↳ As $\beta \uparrow \rightsquigarrow$ KL penalizes stronger \rightsquigarrow latent becomes **statistically i.i.d.**
- ↳ **Statistically i.i.d.** z is **better disentangled**

Posterior Collapse

Posterior Collapse

Posterior collapse refers to the case in which decoder learns to **ignore** latent during **training**

- + What does that mean?!
- It means that the decoder learns to **sample data by its own without caring about encoder!**

Say we are training a VAE: with **posterior collapse** we could say that

- loss computed by **passing encoder's output** through decoder, and
- loss computed by **passing an independent latent** through decoder

are roughly the **same!**

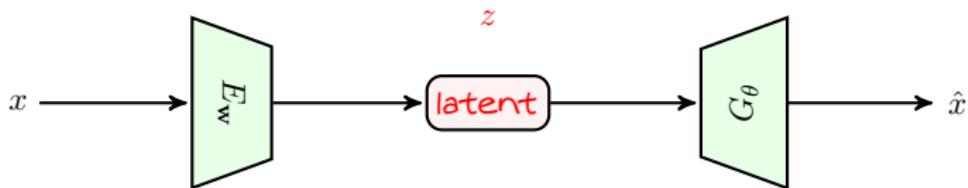
Posterior Collapse

Though posterior collapse does not *directly* impact generation, we can say

- the encoder of such VAE is *useless!*
 - ↳ it gives a *random latent representation* for data samples
 - decoder mainly *learns to fit to the training samples* ↳ when sampling
 - ↳ it can show *worse performance* as it did *not* learn any *structure*
 - ↳ it might *stick to a few samples* ↳ it shows *mode collapse*
- + How do we know if *posterior collapse* is happening?!
- If during training $D_{\text{KL}}(Q_w \parallel \mathcal{N}^0)$ gets *very close to zero*
-
- ③ In VAEs with *powerful decoder*, we can observe *posterior collapse*
 - ↳ We can *control the KL divergence* term as in β -VAE
 - ↳ We can also *reduce* the capacity of the decoder

Treating Sharpness Issue

Looking at sharpness issue: *people thought of old AE-based generation idea*



This old idea says

- ✓ train **deterministic** encoder E_w and decoder G_θ for perfect reconstruction

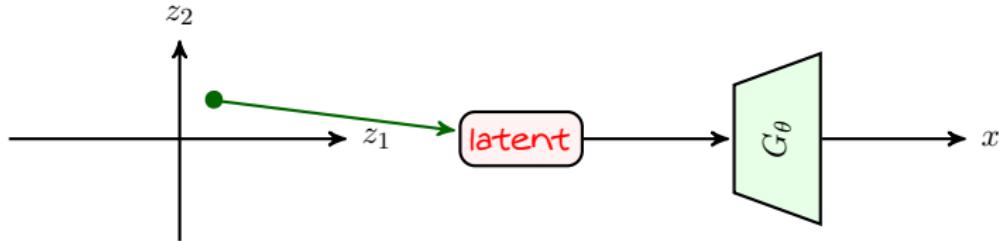
$$R(\mathbf{w}, \theta | \mathbf{x}) = \|\mathbf{x} - G_\theta(E_{\mathbf{w}}(\mathbf{x}))\|$$

- ? learn the distribution of the latent $P(z)$ using encoder
- ✓ sample $z \sim P(z)$ and generate $\mathbf{x} = G_\theta(z)$

Back to AEs: Perfect Reconstruction

- + Why should this basic idea resolves **sharpness issue**!?
- Well! There is no **Gaussian sampling** at the decoder anymore!

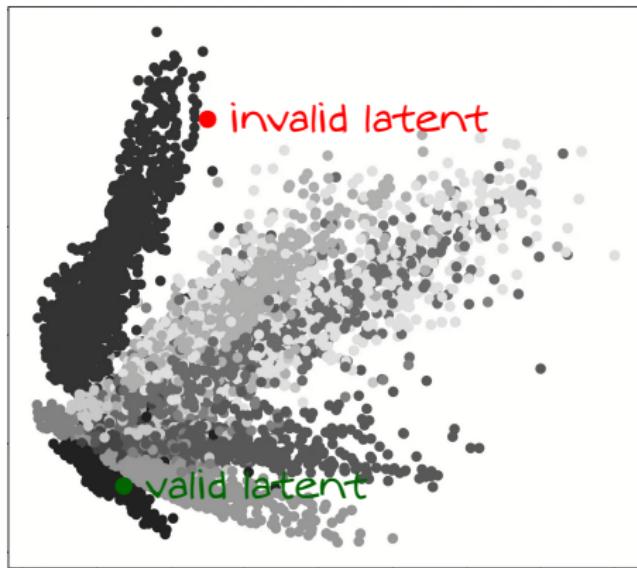
If we are sure z is a **valid latent** \rightsquigarrow we can compute x **perfectly**



Challenge: Unknown Latent Prior

The key challenge is that *in this case*

prior latent distribution is not controlled and hence is unknown!



Learning Latent Prior

- + Can't we **learn** this prior **from the data**!?
- Theoretically, we can!

AE_General_Latent(E_w, G_θ :AE, P_ϕ :Prior_Model):

```
1: Initiate the encoder and decoder
2: for multiple epochs do
3:   Train  $E_w$  and  $G_\theta$  for perfect reconstruction
4: end for
5: for all sample  $x^j$  in dataset do
6:   Compute latent representation  $z^j = E_w(x^j)$ 
7:   Keep it in the latent dataset  $\mathbb{Z} \leftarrow z^j$ 
8: end for
9: for multiple epochs on  $\mathbb{Z}$  do
10:  Train  $P_\phi$                                 #learn latent prior
11: end for
12: return Latent prior and AE
```

Learning Latent Prior

- + What should be P_ϕ ?
- It could be any explicit model we had in Chapter 3

From Chapter 3 we know that

it is not really easy to train P_ϕ on a continuous latent

This motivated for using a discrete latent

- Latent entries belongs to a predefined codebook of k codewords

$$z_i \in \mathcal{C} = \{c_1, \dots, c_k\}$$

- Given that latent space is finite, we can learn P_ϕ more efficiently
 - ↳ For example, we can train an AR model with k -output Softmax at its end

VQ-VAE: Vector Quantized VAE

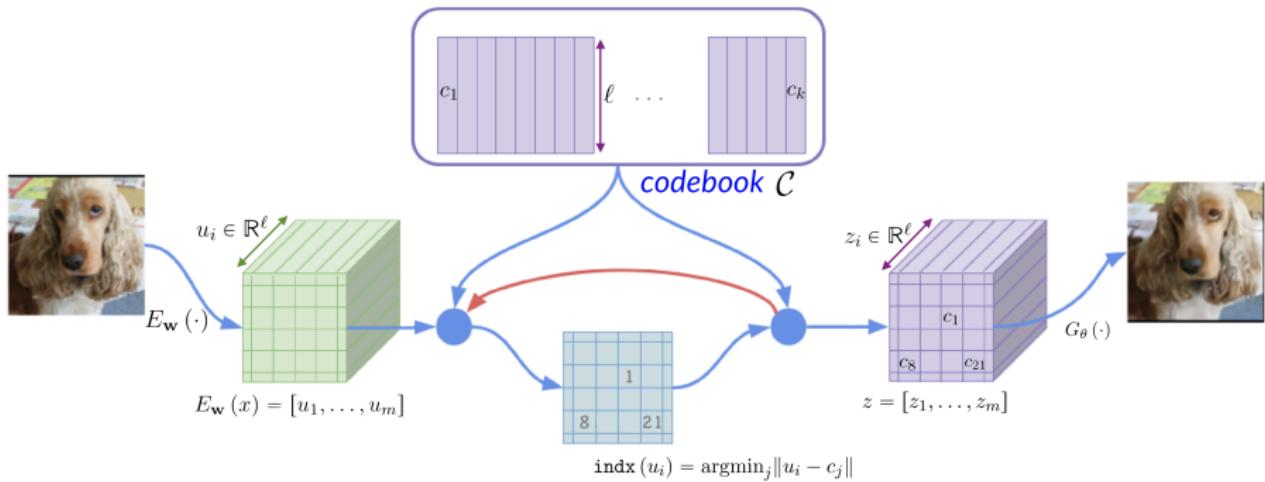
VQ-VAE is an example of generative AEs with *discrete latent*

- ↳ It's usually used for *visual and audio data*
- ↳ It quantizes the outputs of the encoder on a *codebook*
- ↳ The codebook can be seen as an *embedding*
 - ↳ Its codewords, i.e., $c_j \in \mathcal{C}$, are *learnable*

! Attention

Despite its name, it is technically not a **VAE**, since it does **not** use **ELBO** anymore

VQ-VAE: Architecture



Training VQ-VAE

Note that in this architecture E_w , G_θ and codebook are learnable

$$R(\mathbf{w}, \theta, \mathcal{C} | \mathbf{x}) = \|\mathbf{x} - G_\theta(E_w(\mathbf{x}))\|^2 + \beta \|E_w(\mathbf{x}) - z\|^2$$

VQ-VAE: Prior Learning

To learn the **prior**, we invoke AR modeling:

- Define a computational model F_ϕ that extracts **context** up to entry i

$$\mathbf{c}_i = F_\phi(z_{<i}) \in \mathbb{R}^k$$

- Define a conditional distribution as

$$P_\phi(z_i|z_{<i}) = \text{Soft}_{\max}(\mathbf{c}_i) = \text{Soft}_{\max}(F_\phi(z_{<i}))$$

- Model the **latent prior** as

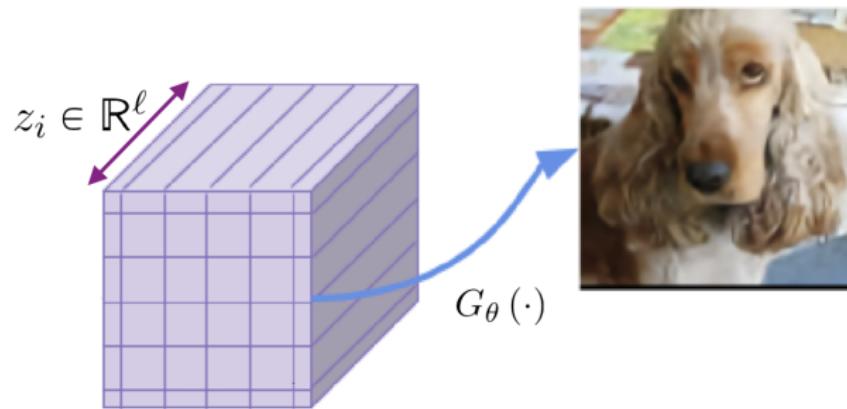
$$P(z) = \prod_i P_\phi(z_i|z_{<i})$$

We train this model on **encoded latent samples** given by the **trained encoder**

VQ-VAE: Sampling

We can sample data distribution by

- ① sampling the AR **latent prior**
- ② passing the **latent** sample through the **decoder**



$$\prod_i P_\phi(z_i | z_{<i})$$

VQ-VAE: Sample Output

In the original VQ-VAE proposal

- CNN encoder and decoder is deployed
- PixelCNN is used to learn latent prior



Wrap Up

In practice VAEs show various favorable properties

- They are very *stable* while training
 - ↳ ELBO maximization \equiv conventional *empirical risk minimization*
- They are *easy* to sample
 - ↳ They are sampled as *flow-based models* with *Gaussian prior*

They could however exhibit

- *blurry output*
 - ↳ This happens due to *Gaussian sampling* at the decoder
- *imperfect disentanglement*
 - ↳ This happens due to *limited control* on the latent
- *posterior collapse*
 - ↳ This happens when the *decoder* is *very strong*

Wrap Up

We can address the **key issues** of VAEs by various **modifications**, e.g., by

- modifying the ELBO term as in β -VAE
- quantizing the **latent** representation as in VQ-VAE

VAEs are still considered as **powerful** generative models

- Like **GANs**, they are **straightforward** to implement in simple applications
- They can provide us further with **latent representation**
 - ↳ This is in particular very useful when we **change modality**

image $\xrightarrow{E_w}$ latent \xrightarrow{LM} text

Next Stop: Indirect distribution learning by **diffusion process**