

Applied Deep Learning

Chapter 6: Recurrent NNs

Ali Bereyhi

`ali.bereyhi@utoronto.ca`

Department of Electrical and Computer Engineering
University of Toronto

Winter 2026

RNN: Dataset and Learning Setting

We are now looking into a supervised learning problem where we are to learn *label* from a *sequence of data* that has generally a *temporal correlation*

Let's denote the *sequence* with $\mathbf{x}[1], \dots, \mathbf{x}[T]$


Temporal Correlation

By *temporal correlation* we mean that entries at *other time instances* carry information about one particular entry $\mathbf{x}[t]$

- + But how is a *label* assigned to this sequence?
- Well, that can be of various forms!

Types of Problems with Sequence Data

We considered a very simple case: *many-to-many type I*




In this case, every entry has a *label*

↳ *speech tagging*: $x[t]$ is a part of speech and $y[t]$ is its tag

Types of Problems with Sequence Data

We considered a very simple case: *many-to-many type II*




In this case, we get a *label* once after multiple entries

↳ *speech recognition: $x[t]$ is a small part of speech and $y[t]$ says what is a every couple of minutes about*

Types of Problems with Sequence Data

We considered a very simple case: *many-to-many type III*




In this case, we start to get *labels* after some delay

↳ language translation: $x[1], \dots, x[t-1]$ is a sentence in German and $y[t-1], \dots, y[T]$ is its translation to English

Types of Problems with Sequence Data

We considered a very simple case: *one-to-many*




In this case, we get only one input data and have a sequence of *labels*

↳ image captioning: $x[1]$ is an image and $y[1], \dots, y[T]$ is a caption describing what is inside the image

Types of Problems with Sequence Data


We considered a very simple case: *many-to-one*



In this case, we get only one *label* for a whole input sequence

↳ sequence classification: $x[1], \dots, x[T]$ is a speech and $y[T]$ says if this speech is constructive or destructive

Types of Problems with Sequence Data: *FNNs*




In fact, FNNs are *one-to-one* RNNs

- ↳ we can think of every data-point as *a sequence of length one*, or
- ↳ we may think of dataset as a long sequence with *no temporal correlation*

Shallow RNN

Let's break it down a bit: say the **yellow box** is a **fully-connected layer**




This layer gets input $\mathbf{x} \in \mathbb{R}^N$ and returns **activated** feature $\mathbf{h} \in \mathbb{R}^M$

- We replace its **activation** by $\tanh(\cdot)$
 - ↳ Not necessary, but usually suggested
- We modify it to get a new input in $\mathbb{R}^{N+M}$
 - ↳ N entries for **data inputs** $\mathbf{x}$ in time t
 - ↳ M entries for **features** $\mathbf{h}$ in time $t - 1$

Shallow RNN

Let's break it down a bit: say the **yellow box** is a **fully-connected** layer




We show it this way now

- We call $h[t]$ usually the **hidden state**
- We pass the **hidden state** through an output layer
 - ↳ Output is **not necessarily** corresponding to label

Shallow RNN

We may show our **shallow RNN** **compactly** via the following diagram




In this diagram

- Each edge is a set of weights, e.g., a weight matrix
 - The return edge also has a delay in time
- + But, isn't that simply **Elman Network**?
- If we use a fully-connected layer and sigmoid activation; then, Yes! But, Remember that
 - Elman did **not** **train it over time**
 - Elman in its model used **sigmoid activation**

Deep RNN

We can add more layers to make a **deep** RNN




In this diagram

- Each edge is a set of weights, e.g., a weight matrix
- The return edge also has a delay in time
- And, it's no more Elman Network

Deep RNN

It might be easier to think of it as the following diagram



We can use any module that we like

- We may add *skip connection*
- We could make *dense connections*
- We may *skip over time*
- We may replace hidden layers with *convolutions* or anything else


Shallow RNN: *Elman-Like Network*

Let's start with simple RNN: a shallow RNN with *fully-connected* hidden layer

+ You mean *Elman Network*?


- Yes, but we now try to address the *challenges* that Elman did not address

Let's look at the flow of information once again



Shallow RNN: Forward Propagation

Let's specify the learnable parameters with some colors



Say we set the activation to $f(\cdot)$

- ① We have $y[t] = f(\mathbf{W}_2 \mathbf{h}[t])$: we *can learn* $\mathbf{W}_2$
- ② We have $\mathbf{h}[t] = f(\mathbf{W}_1 \mathbf{x}[t] + \mathbf{W}_m \mathbf{h}[t-1])$: we *can learn* $\mathbf{W}_1$ and $\mathbf{W}_m$
- ③ We can start with any *hidden state*: this means we *can learn* $\mathbf{h}[0]$ too